

Ankur Sarda

(518) 466-8710 | sardankur@gmail.com

SUMMARY

Staff Software Engineer with 10+ years of experience building AI applications, ML infrastructure, and Data systems. Led platform initiatives from concept to production, including LLM evaluation frameworks, anomaly and fraud detection systems, distributed inference services, and graph platforms. Experience across Oracle, Uber, and Goldman Sachs.

EXPERIENCE

Software Engineer | Oracle

2021 - Present

- Led development of Veritas, a distributed LLM evaluation platform supporting 150+ evaluation recipes across RAG, NL2SQL, Responsible AI, reasoning, and multimodal use cases; built Ray-based multi-node execution to orchestrate mixed IO- and CPU-heavy workloads, including network model calls and local model/tool execution; delivered 50+ model evaluation reports for OpenAI, Anthropic, Meta, and Cohere.
- Standardized LLM evaluation workflows across organizations, introducing leaderboard-driven model comparison and reducing ad-hoc evaluation effort while improving consistency in benchmarking.
- Drove development of three SLURM-based GPU training clusters, each with 5K+ GPUs, consolidating distributed compute resources across Oracle and increasing utilization from ~30% to 80%+.
- Scaled an effort to curate datasets for training and evaluation, including 100K de-identified clinical doctor-patient records with assisted PII/PHI removal, obfuscation, and human annotation, plus 10K NL2SQL benchmarks generated from public databases with systematic coverage across joins, filters, aggregations, GROUP BY, CTEs, and table complexity.
- Led cross-functional development of AI Asset Inventory, an enterprise AI governance platform for tracking AI systems across Oracle. Drove the product from design through GA across multiple organizations and integrated it with Oracle AI Review, a mandatory productionization workflow.
- Built a distributed cost-aware anomaly detection inference service, enabling efficient model execution and routing based on capacity across diverse workloads reducing latency by 60% and improving compute utilization by 30%.
- Contributed to company-wide hiring as a bar-raiser interviewer (150+ interviews)

Software Engineer | Uber

2018 - 2021

- Built and productionized large-scale graph learning systems for Uber Eats recommendations and fraud detection, processing billions of entities across users, devices, vehicles, and restaurants. Curated graph data using Cypher over Spark and piped together workflows for weekly training, daily prediction, and hyperparameter optimization. Work was featured in Uber Engineering blog posts on [Uber Fraud Detection](#) & [Uber Eats graph learning](#).
- Led cross-team migration of 100+ critical risk assessment metrics to Apache Flink-based streaming infrastructure, reducing detection latency from 10+ minutes to under 10 seconds while building consensus, driving adoption, and deprecating legacy metric pipelines.
- Designed and contributed to a next-generation rider identity verification platform, leveraging social and government identity signals to improve risk assessment and fraud prevention.
- Contributed to onboarding and training programs for new hires as an instructor for Uber's data ecosystem.
- As an intern, developed an Ads attribution pipeline and dashboard for Uber's Ad expenditure profile, providing insights on growth, expense, ROI etc.

Earlier Experience

Previously worked on Spark and Scala-based data systems at Goldman Sachs, building distributed data pipelines. Also interned at Intuit developing Android applications and at McAfee on antivirus products.

EDUCATION

Stony Brook University, New York — M.S. in Computer Science, 2017 | GPA 3.91/4.0

PES Institute of Technology, Bangalore — B.E. in Computer Science, 2014 | GPA 9.29/10